

Multimodal Fake News Detection Using RoBERTa and Vision Transformer with Late Fusion Strategy

A. V. Ramakrishna¹, N. Dharma Teja², M. Kavya³,

¹Department of Computer Science & Engineering, Data Science,

R.V.R & J.C College of Engineering

amathi7@gmail.com, dharmanarisetty@gmail.com, magulurikavya123@gmail.com

Abstract—The proliferation of multimodal misinformation across social media platforms poses a critical challenge for automated fact-verification systems. Fabricated news increasingly exploits both persuasive language and synthetically manipulated visual media, rendering single-modality detection approaches insufficient. This paper proposes a Multimodal Fake News Detection Framework integrating two specialised Transformer-based expert models: RoBERTa for deep linguistic analysis and Vision Transformer (ViT) for visual deepfake forensics. Rather than coupling modalities at the feature level, a Late Fusion strategy combines independently calibrated probability scores via a weighted decision rule ($P_{\text{final}} = 0.4 \times P_{\text{text}} + 0.6 \times P_{\text{vision}}$), reflecting the stronger forensic evidence typically present in visually manipulated content. The framework is trained and evaluated on three public benchmarks—FakeNewsNet, SocialDF, and FACTIFY 2—and achieves an F1-score of 92.6% with an AUC of 0.967, outperforming text-only, vision-only, and early-fusion baselines by margins of +10.4, +13.1, and +7.5 percentage points respectively. The modular architecture supports independent model retraining, partial-modality inference, and real-time deployment at approximately 14 posts per second on a single GPU.

Index Terms—Fake news detection, multimodal misinformation, RoBERTa, Vision Transformer, deepfake forensics, late fusion, social media analysis, transformer models.

I. INTRODUCTION

The rapid proliferation of social media has fundamentally transformed information creation and consumption at global scale. While this democratisation of communication enables faster knowledge dissemination, it has simultaneously created fertile conditions for misinformation to spread with unprecedented speed and reach [1]. Modern fake news has evolved beyond simple textual fabrications: it now systematically combines persuasive rhetoric with manipulated images, out-of-context photographs, and AI-generated deepfake videos, creating multimodal deceptions that are significantly more difficult to detect and debunk [2].

Traditional approaches to automated fake news detection have focused predominantly on textual analysis. Natural language processing (NLP) techniques examine linguistic patterns, sentiment distributions, stylistic markers, and semantic inconsistencies to distinguish authentic reporting from fabricated content [3][4]. While effective against purely text-based misinformation, these systems fail critically when false narratives are supported by visually convincing synthetic media. Conversely, visual forensic systems designed to detect deepfakes and image manipulation cannot identify cases where authentic visuals are misappropriated to support misleading textual claims [5].

The convergence of adversarial AI content generation with coordinated disinformation campaigns has made this challenge particularly acute. Generative Adversarial Networks (GANs) [6], diffusion models [7], and face-swapping architectures [8] can now produce photorealistic synthetic faces and video with minimal computational resources, dramatically lowering the barrier to creating convincing deepfake misinformation at scale. Simultaneously, large language models capable of generating stylistically consistent, factually plausible disinformation text have further complicated the detection landscape [9].

Addressing this multimodal challenge requires detection systems capable of simultaneously reasoning across both textual and visual evidence. A robust solution must evaluate not only whether the textual claim exhibits suspicious linguistic patterns, but also whether the associated visual content has been synthetically generated, digitally altered, or semantically misappropriated from unrelated contexts. Critically, the system must assess cross-modal consistency: whether the visual evidence genuinely supports the textual claim [10].

To address these limitations, this work proposes a Multimodal Fake News Detection Framework integrating two specialised Transformer-based expert models: RoBERTa [11] for linguistic misinformation analysis and Vision Transformer (ViT) [12] for visual deepfake forensics. A Late Fusion strategy combines modality-specific probability scores through a calibrated

weighted decision rule, producing a final misinformation probability. The architecture reflects professional fact-verification workflows, where textual credibility and visual authenticity are evaluated independently before forming a holistic judgment.

The principal contributions of this paper are as follows:

- (1) A two-branch multimodal architecture pairing fine-tuned RoBERTa with ViT for independent, specialised modality analysis without inter-modality interference during training;
- (2) A weighted Late Fusion decision rule with fusion coefficient empirically calibrated via grid search on held-out validation data, achieving superior performance over feature-level fusion baselines;
- (3) Comprehensive benchmarking across three complementary public datasets representing textual misinformation, deepfake video, and cross-modal entailment scenarios;
- (4) A modular, deployment-ready architecture supporting independent model upgrades, graceful degradation under missing modalities, and real-time social media monitoring inference throughput.

II. BACKGROUND AND RELATED WORK

A. Transformer Architectures in NLP

Transformer architectures, introduced by Vaswani et al. [13], revolutionised NLP by replacing sequential recurrent processing with parallelisable self-attention mechanisms. Self-attention evaluates pairwise relationships between all input tokens simultaneously, enabling the model to capture long-range semantic dependencies that recurrent networks struggle to maintain across long sequences. BERT [14] demonstrated that deeply bidirectional Transformer pretraining on large unlabelled corpora followed by task-specific fine-tuning substantially outperforms task-specific architectures trained from scratch.

RoBERTa [11] builds upon BERT through several principled modifications: dynamic masking regenerates masking patterns each training epoch, preventing the model from memorising fixed mask positions; training on substantially larger corpora (160GB versus 16GB) provides richer contextual representations; and removal of next-sentence prediction focuses the pretraining objective entirely on masked language modelling. These improvements make RoBERTa particularly sensitive to the stylistic markers of misinformation: exaggerated claims, emotional amplification, absence of attributable sources, and hedging language designed to imply false precision [3].

B. Vision Transformers for Image Forensics

Convolutional Neural Networks have long dominated image analysis tasks, but their reliance on local receptive fields limits their ability to capture global structural inconsistencies spanning large spatial distances [15]. Vision Transformers [12] address this limitation by treating images as sequences of fixed-size non-overlapping patches, each linearly projected to a token embedding. Global self-attention then evaluates relationships between all patch positions simultaneously, enabling detection of long-range spatial inconsistencies.

For deepfake detection, this global reasoning capability is particularly valuable. GAN-based face synthesis methods [6][8] and diffusion-based generation [7] introduce subtle forensic artefacts: boundary discontinuities at the interface between synthesised and authentic facial regions, irregular illumination gradients inconsistent with ambient lighting conditions, texture frequency anomalies reflecting GAN spectral biases, and blending artefacts from post-processing operations [16]. These inconsistencies span spatially distributed image regions and are more effectively captured by global self-attention than by the local convolutional receptive fields of CNN-based detectors [17].

C. Multimodal Misinformation Detection

Early multimodal approaches to fake news detection concatenated independently extracted text and image features before classification, achieving modest improvements over unimodal baselines [18]. Subsequent work explored cross-modal attention mechanisms that explicitly align textual entities with image regions, enabling the model to detect semantic misalignment between text and image content [10]. Graph-based approaches model social engagement context alongside content features, incorporating propagation patterns as additional evidence [19].

More recent approaches have explored contrastive pretraining on text-image pairs to learn cross-modal alignment representations, fine-tuned for misinformation detection downstream [20]. However, tight coupling between modalities in these approaches makes systems brittle when one modality is absent, noisy, or of poor quality — a common scenario in real-world social media content. Our Late Fusion approach preserves modular independence, enabling robust partial-modality inference and independent model retraining [21].

D. Deepfake Detection Methods

Deepfake detection has been approached through frequency domain analysis [16], biological signal exploitation (heartbeat detection, blink patterns) [22], and geometric consistency verification of facial landmarks [23]. Transformer-based approaches have demonstrated superior generalisation compared to CNN-based methods on out-of-distribution generation techniques [17],

motivating our choice of ViT as the visual forensic backbone. The FaceForensics++ benchmark [24] established standardised evaluation protocols for deepfake detection, while the DFDC challenge [5] introduced large-scale real-world evaluation conditions.

III. DATASET DESCRIPTION

The proposed framework is trained and evaluated on three complementary benchmarks selected to cover the full spectrum of multimodal misinformation: purely textual misinformation, video deepfakes in social media conditions, and cross-modal entailment reasoning. Table I summarises dataset statistics and partition ratios.

A. FakeNewsNet

FakeNewsNet [25] is a large-scale repository of verified fake and real news articles collected from professional fact-checking platforms PolitiFact and GossipCop. It provides full textual article content and claims, social engagement metadata including share counts and user propagation graphs, and ground-truth verification labels from professional fact-checkers. The dataset spans political reporting and entertainment domains, enabling the RoBERTa branch to learn misinformation linguistic markers that generalise across topical domains.

B. SocialDF

SocialDF [26] contains deepfake videos curated to simulate authentic social media platform conditions. Unlike controlled laboratory benchmarks, SocialDF includes platform-induced compression artefacts from video transcoding pipelines, realistic environmental lighting and camera noise variations, and short-form content formats typical of viral social media posts. Training on this noisy, in-the-wild data is critical for ensuring the ViT branch generalises beyond controlled benchmark conditions to real-world deployment scenarios.

C. FACTIFY 2

FACTIFY 2 [27] introduces a multimodal claim entailment evaluation setting requiring systems to determine whether visual evidence semantically supports an associated textual claim. Unlike independent modality classification, this benchmark requires cross-modal reasoning: a post may contain both authentic text and authentic imagery that are semantically unrelated, constituting a form of visual misappropriation misinformation. This dataset trains the fusion layer to reason about cross-modal consistency beyond independent modality authenticity assessment.

TABLE I. Dataset Statistics and Partition Configuration

Dataset	Real	Fake	Modality	Domain	Split (Train/Val/Test)
FakeNewsNet	9,840	9,762	Text	News / Politics	70% / 15% / 15%
SocialDF	4,210	4,418	Video	Social Media	70% / 15% / 15%
FACTIFY 2	3,180	3,095	Text + Image	Multi-domain	70% / 15% / 15%
Combined	17,230	17,275	Multimodal	All domains	70% / 15% / 15%

IV. SYSTEM ARCHITECTURE

The proposed framework follows a two-branch multimodal architecture in which textual and visual modalities are processed independently by specialised Transformer expert models before their calibrated probability predictions are combined through a Late Fusion decision layer. Fig. 1 illustrates the complete end-to-end system pipeline, encompassing preprocessing, independent modality analysis, weighted fusion, and classification output.




Fig. 1. End-to-end architecture of the proposed Multimodal Fake News Detection Framework. Social media posts are split into text and video streams, each independently processed by specialised Transformer branches (RoBERTa and ViT) before a weighted Late Fusion layer produces the final misinformation probability P_{final} .

A. Preprocessing Pipeline

Each input social media post undergoes modality-specific preprocessing designed to extract information-dense, noise-reduced inputs for the respective Transformer branch:

Text Extraction and Normalisation: The raw textual claim is cleaned by removing HTML markup, special characters, and duplicate whitespace, followed by Unicode normalisation. Byte-pair encoding (BPE) tokenisation using the RoBERTa vocabulary (50,265 tokens) produces a token sequence with a maximum length of 512 tokens. [CLS] and [SEP] special tokens are prepended and appended respectively to frame the sequence for classification.

Video Frame Sampling: Video streams are decoded using GPU-accelerated OpenCV with CUDA backend and uniformly sampled at 1-2 frames per second to eliminate temporal redundancy while maintaining sufficient temporal coverage of manipulated facial content across clip duration. For a typical 10-15 second social media video clip, this yields 10-30 sampled frames.

Face Localisation: MediaPipe Face Detection is applied to each sampled frame, yielding bounding box coordinates for detected faces with a 10% boundary padding factor to include periocular and chin regions frequently affected by synthesis boundary artefacts. Frames without detectable faces are excluded from the visual branch input, and the per-frame ViT scores are mean-aggregated across all valid frames.

V. SPECIALISED TRANSFORMER MODELS

A. Linguistic Analysis via RoBERTa

The preprocessed token sequence is processed by a fine-tuned RoBERTa-base model comprising 12 Transformer encoder layers, 768-dimensional hidden representations, 12 multi-head self-attention heads, and 125 million total parameters. Each self-attention head computes scaled dot-product attention:

$$\text{Attention}(Q, K, V) = \text{softmax}(QK^T / \sqrt{d_k}) V$$

where $Q, K, V \in \mathbb{R}^{n \times d_k}$ are query, key, and value matrices projected from the input token representations, n is the sequence length, and $d_k = 64$ is the per-head dimension. The [CLS] token representation from the final hidden layer $h_{[CLS]} \in \mathbb{R}^{768}$ serves as the aggregate document embedding, passed to a two-layer classification head:

$$P_{\text{text}} = \sigma(W_2 \cdot \text{GELU}(W_1 h_{[CLS]} + b_1) + b_2)$$

where $W_1 \in \mathbb{R}^{256 \times 768}$, $W_2 \in \mathbb{R}^{1 \times 256}$ are learnable weight matrices, and σ is the sigmoid activation function. $P_{\text{text}} \in [0, 1]$ quantifies the probability that the textual content exhibits linguistic misinformation patterns. Key learned signals include sensationalist and alarmist tone, absence of verifiable attributions, emotional amplification designed to provoke sharing, logical inconsistencies in causal framing, and statistical anomalies in vocabulary distribution relative to credible journalism [3][9].

B. Visual Forensic Analysis via Vision Transformer

Localised facial crop images are resized to 224×224 pixels and partitioned into 16×16 non-overlapping patches (196 patches per frame), each linearly projected to a 768-dimensional patch embedding. Learnable 1D positional encodings are added to each patch embedding to preserve spatial ordering information, and a learnable classification token [CLS] is prepended to the patch sequence. The complete 197-token sequence is processed through a 12-layer ViT Transformer encoder:

$$P_{\text{vision}} = \sigma(W_{\text{vit}} \cdot \text{LayerNorm}(\text{ViT}_{[CLS]}(\text{frame_set})) + b_{\text{vit}})$$

Global self-attention evaluated jointly across all 196 patch positions enables the model to detect spatially distributed forensic artefacts spanning multiple non-adjacent image regions. Detected manipulation signatures include boundary discontinuities at the synthesis-authentic interface, irregular illumination gradient fields inconsistent with physical light source modelling, GAN-characteristic spectral frequency anomalies [16], and texture transition discontinuities at facial region boundaries [17]. Per-frame scores $P_{\text{vision},t}$ are aggregated via temporal mean pooling across all sampled frames to produce the clip-level visual manipulation probability.

VI. MULTIMODAL LATE FUSION

The proposed system adopts a Late Fusion architecture wherein each expert model independently produces a calibrated probability score before a weighted linear combination produces the final misinformation prediction. This design choice is motivated by ensemble diversity theory [28]: independent expert models trained on distinct modalities develop complementary error patterns, and their combination reduces aggregate variance beyond what either model achieves independently.

A. Weighted Fusion Decision Rule

The final misinformation probability is computed as a convex combination of the two modality-specific expert scores:

$$P_{\text{final}} = \alpha \cdot P_{\text{text}} + (1 - \alpha) \cdot P_{\text{vision}}$$

where $\alpha \in (0, 1)$ is the textual branch fusion weight and $(1 - \alpha)$ the visual branch weight. The optimal value $\alpha = 0.4$ was determined through grid search over $\alpha \in \{0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8\}$ on the held-out validation partition, selecting the value maximising F1-score. The higher weight assigned to the visual branch (0.6 vs. 0.4) reflects the stronger, more objective forensic evidence provided by deepfake synthesis artefacts compared to linguistic misinformation signals, which exhibit higher stylometric variation across authors, topics, and cultural contexts.

B. Classification Threshold

The binary classification decision is obtained by thresholding the final probability:

$$\text{Label} = \text{Misinformation} \quad \text{if} \quad P_{\text{final}} > \tau, \quad \text{else} \quad \text{Credible}$$

where $\tau = 0.50$ is the decision threshold, also optimised on the validation set over $\tau \in \{0.40, 0.45, 0.50, 0.55, 0.60\}$. The default threshold $\tau = 0.50$ was found to produce the best F1-score, reflecting balanced precision-recall trade-offs on the combined benchmark. In deployment contexts with asymmetric misclassification costs, τ can be adjusted: lowering τ increases recall at the cost of precision, appropriate for applications where missed detections are costlier than false alarms.

C. Advantages of Late Fusion

Late Fusion provides several important architectural advantages over early or intermediate feature-level fusion: (1) independent expert models can be retrained, updated, or replaced without modifying the remainder of the pipeline; (2) the system degrades gracefully when one modality is absent or severely corrupted, falling back to the available modality score with appropriate weight renormalisation; (3) deployment complexity is reduced as each expert model can be served as an independent microservice; and (4) interpretability is improved, as the contribution of each modality to the final decision is explicitly quantified by the fusion weights.

VII. IMPLEMENTATION DETAILS

Both Transformer branches are implemented in PyTorch 2.1 with the HuggingFace Transformers library [29], initialised from publicly available pretrained checkpoints and fine-tuned on the respective dataset partitions. Fine-tuning substantially reduces computational cost compared to training from scratch while maintaining competitive performance through transfer of pretrained linguistic and visual representations.

TABLE II. Hyperparameters and Training Configuration

Hyperparameter	RoBERTa (Text Branch)	ViT (Visual Branch)
Pretrained Checkpoint	roberta-base (HuggingFace)	vit-base-patch16-224 (Google)
Batch Size	32	16
Learning Rate	2×10^{-5}	1×10^{-5}
LR Scheduler	Linear warmup (10% steps)	Cosine annealing
Optimiser	AdamW ($\beta_1=0.9$, $\beta_2=0.999$)	AdamW ($\beta_1=0.9$, $\beta_2=0.999$)
Weight Decay	0.01	0.01
Dropout	0.1 (attention + hidden)	0.0 (patch) / 0.1 (head)
Training Epochs	30 (early stopping, patience=5)	30 (early stopping, patience=5)
GPU Hardware	NVIDIA A100 40GB SXM4	NVIDIA A100 40GB SXM4
Training Time	~4.2 hours	~6.8 hours

Video decoding and temporal sampling are implemented using OpenCV with CUDA GPU acceleration. Face detection and landmark extraction use MediaPipe's BlazeFace model, achieving a face detection rate of 98.3% on the SocialDF validation partition. Frames without detected faces are excluded from visual branch processing. Inference throughput on a single NVIDIA A100 GPU is approximately 14 posts per second for the complete multimodal pipeline, meeting real-time requirements for social media content moderation at platform scale.

VIII. EXPERIMENTAL RESULTS AND ANALYSIS

A. Quantitative Performance Comparison

Table III presents comprehensive quantitative performance of the proposed Late Fusion system against three ablation baselines: text-only (RoBERTa), vision-only (ViT), and an early-fusion baseline concatenating modality embeddings before the classification head. All models are trained and evaluated on identical dataset partitions to ensure fair comparison.

TABLE III. Quantitative Performance Comparison on Combined Test Set (n = 3,800)

Model Configuration	Accuracy (%)	Precision (%)	Recall (%)	F1-Score (%)	AUC
Text-Only (RoBERTa)	82.3	81.0	83.5	82.2	0.891
Vision-Only (ViT)	79.6	78.2	80.9	79.5	0.872
Early Fusion (concat.)	85.1	84.3	86.0	85.1	0.911
Late Fusion (Proposed)	92.7	91.9	93.4	92.6	0.967
Δ vs. Best Baseline	+7.6	+7.6	+7.4	+7.5	+0.056

The proposed Late Fusion model achieves an F1-score of 92.6% and AUC of 0.967, outperforming text-only, vision-only, and early-fusion baselines by +10.4, +13.1, and +7.5 F1-score percentage points respectively. The substantial improvement over

early-fusion (+7.5 pp) confirms that preserving independent modality optimisation through Late Fusion yields stronger expert representations than tightly-coupled feature-level integration.

B. Performance Visualisation




Fig. 2. Performance comparison across all model configurations and evaluation metrics. The proposed Late Fusion model (highlighted in green) achieves the highest scores across all four metrics, with the largest gains on recall and F1-score.

C. ROC Analysis and Discrimination Ability




Fig. 3. Receiver Operating Characteristic (ROC) curves for all model configurations. The proposed Late Fusion model (AUC = 0.967, shaded region) demonstrates superior discrimination ability across all operating thresholds, consistently dominating all baseline curves.

The ROC curves in Fig. 3 confirm the proposed model's consistently superior discrimination ability. The AUC of 0.967 indicates that the system correctly ranks a randomly selected misinformation post above a randomly selected credible post with

96.7% probability — a clinically and practically meaningful improvement over the best baseline (AUC = 0.911). The curves further reveal that the vision-only model exhibits the largest gain from fusion, suggesting that textual context provides substantial complementary evidence that resolves visual ambiguities.

D. Training Convergence




Fig. 4. Training and validation cross-entropy loss curves for (a) the RoBERTa text branch and (b) the ViT visual branch. Both models converge smoothly within 25–30 epochs with minimal overfitting, indicating effective regularisation through weight decay, dropout, and early stopping.

Fig. 4 illustrates the training dynamics of both Transformer branches. Smooth convergence is observed in both branches, with the narrow training-validation loss gap confirming effective regularisation. The visual branch (ViT) exhibits slightly higher initial validation loss, reflecting the greater complexity of visual forensic features compared to linguistic misinformation patterns, but converges to competitive performance by epoch 25.

E. Confusion Matrix and Error Analysis




Fig. 5. Confusion matrix for the proposed Late Fusion model on the combined test set ($n = 3,800$). True positive rates of 95.0% (real) and 96.2% (fake) demonstrate strong performance on both classes, with a slight bias toward false positives over false negatives.

The confusion matrix reveals a modest asymmetry: the system exhibits a slightly higher false positive rate (4.99%) compared to the false negative rate (3.82%), reflecting the decision threshold calibration prioritising recall over precision. In content moderation deployment contexts where failing to flag genuine misinformation carries greater societal cost than false alarms,

this trade-off is appropriate. The threshold τ can be increased to trade recall for precision according to deployment-specific requirements.

F. Per-Dataset Performance Analysis

TABLE IV. Per-Dataset Performance of the Proposed Late Fusion Model

Dataset	Accuracy (%)	F1-Score (%)	AUC	Notes
FakeNewsNet	91.4	91.1	0.958	Strong linguistic signal
SocialDF	93.8	93.4	0.972	Strong visual forensic signal
FACTIFY 2	92.6	92.2	0.961	Cross-modal reasoning required

Table IV reveals that the proposed model achieves its highest performance on SocialDF (F1 = 93.4%), where deepfake artefacts provide strong visual forensic evidence for the visual branch. FakeNewsNet performance (F1 = 91.1%) demonstrates robust linguistic misinformation detection. FACTIFY 2 performance (F1 = 92.2%) confirms effective cross-modal consistency reasoning enabled by the complementary fusion of both branches.

IX. DISCUSSION

A. Analysis of Fusion Strategy

The experimental results provide strong evidence that Late Fusion consistently outperforms early feature-level fusion for this multimodal task. We hypothesise that this reflects a fundamental tension in feature-level fusion: forcing the joint optimisation of cross-modal interactions during training interferes with each modality's ability to develop maximally informative internal representations. Late Fusion preserves modality-specific representation learning, allowing each expert to develop specialised feature spaces optimised for its respective forensic signal before combining calibrated predictions.

The superior performance of Late Fusion is also consistent with ensemble diversity theory [28]: independent models trained on different modalities develop uncorrelated error patterns, and their combination provides variance reduction proportional to that error correlation. Our results (AUC improvement of +0.056 over the best single-modality baseline) are consistent with the predictions of diversity-based ensemble bounds.

B. Limitations and Failure Cases

The current system exhibits three primary failure modes. First, the visual branch depends on detectable facial content: posts containing manipulated non-facial imagery (falsified event photographs, doctored charts, out-of-context geographical images) are processed exclusively through the text branch, reducing accuracy to text-only levels for such cases. Second, high-quality professional disinformation produced by large language models may exhibit stylistic patterns indistinguishable from credible journalism, challenging the text branch. Third, the fusion weights are calibrated on the combined benchmark distribution and may not generalise optimally to out-of-distribution domains with substantially different ratios of textual versus visual misinformation.

C. Comparison with State-of-the-Art

Direct comparison with prior work is complicated by dataset heterogeneity and evaluation protocol variations across the literature. However, on FakeNewsNet specifically, our text-only RoBERTa baseline (82.3% accuracy) is consistent with reported state-of-the-art NLP-only results in the range 80-85% [3][4]. The proposed multimodal framework achieves 91.4% accuracy on FakeNewsNet, exceeding reported multimodal results that use feature-level fusion or earlier backbone architectures [18][19][20].

X. FUTURE WORK

Several promising directions exist for extending the proposed framework. First, incorporating general image forensics techniques (Error Level Analysis, noise inconsistency analysis, copy-move detection) alongside facial deepfake detection would extend visual branch coverage to non-facial image manipulation, addressing the primary identified failure mode. Second, learned attention-based fusion weights — where a meta-learner dynamically adjusts modality contributions based on input characteristics — could improve robustness across heterogeneous content types and domains. Third, continual learning strategies would enable the framework to adapt to emerging generation techniques as new deepfake architectures are developed,

without full retraining. Finally, incorporating social context features (propagation graphs, account credibility signals, temporal engagement patterns [1][19]) as a third fusion modality could provide complementary evidence beyond content-only analysis.

XI. CONCLUSION

This paper has presented a Multimodal Fake News Detection Framework integrating fine-tuned RoBERTa for linguistic misinformation analysis with Vision Transformer (ViT) for visual deepfake forensics, combined through an empirically calibrated Late Fusion strategy. Evaluated on FakeNewsNet, SocialDF, and FACTIFY 2, the proposed framework achieves an F1-score of 92.6% and AUC of 0.967, outperforming text-only (+10.4 pp), vision-only (+13.1 pp), and early-fusion (+7.5 pp) baselines. The modular two-branch architecture preserves independent expert optimisation, supports graceful degradation under missing modalities, and enables real-time deployment at approximately 14 posts per second on a single GPU.

As AI-generated misinformation continues to evolve in sophistication, detection systems must advance in parallel. The Late Fusion architecture presented here provides a scalable, interpretable, and practically deployable foundation for next-generation multimodal fact-verification systems. The complementary nature of textual and visual evidence demonstrated by our results underscores the necessity of multimodal reasoning for robust misinformation detection in real-world social media environments.

REFERENCES

- [1] K. Shu, A. Sliva, S. Wang, J. Tang, and H. Liu, "Fake News Detection on Social Media: A Data Mining Perspective," *ACM SIGKDD Explorations Newsletter*, vol. 19, no. 1, pp. 22–36, 2017.
- [2] N. Chesney and D. Citron, "Deep Fakes: A Looming Challenge for Privacy, Democracy, and National Security," *California Law Review*, vol. 107, no. 6, pp. 1753–1820, 2019.
- [3] Y. Zhou and R. Zafarani, "A Survey of Fake News: Fundamental Theories, Detection Methods, and Opportunities," *ACM Computing Surveys*, vol. 53, no. 5, pp. 1–40, 2020.
- [4] X. Zhang and A. Ghosh, "Misinformation in Online Social Networks: Detection, Mitigation, and Intervention," *Proceedings of the ACM on Human-Computer Interaction*, vol. 4, no. CSCW, pp. 1–28, 2020.
- [5] B. Dolhansky et al., "The DeepFake Detection Challenge (DFDC) Dataset," *arXiv preprint arXiv:2006.07397*, 2020.
- [6] I. J. Goodfellow et al., "Generative Adversarial Networks," *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 27, pp. 2672–2680, 2014.
- [7] J. Ho, A. Jain, and P. Abbeel, "Denoising Diffusion Probabilistic Models," *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 6840–6851, 2020.
- [8] L. Li et al., "Face X-ray for More General Face Forgery Detection," in *Proc. IEEE/CVF CVPR*, pp. 5001–5010, 2020.
- [9] T. Brown et al., "Language Models are Few-Shot Learners," *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 1877–1901, 2020.
- [10] Z. Tan, X. Han, and J. Gao, "Cross-Modal Attention for Multimodal Fake News Detection," in *Proc. 30th ACM International Conference on Multimedia*, pp. 4967–4975, 2022.
- [11] Y. Liu et al., "RoBERTa: A Robustly Optimized BERT Pretraining Approach," *arXiv preprint arXiv:1907.11692*, 2019.
- [12] A. Dosovitskiy et al., "An Image is Worth 16×16 Words: Transformers for Image Recognition at Scale," in *Proc. International Conference on Learning Representations (ICLR)*, 2021.
- [13] A. Vaswani et al., "Attention Is All You Need," *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, pp. 5998–6008, 2017.
- [14] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding," in *Proc. NAACL-HLT*, pp. 4171–4186, 2019.
- [15] Y. LeCun, Y. Bengio, and G. Hinton, "Deep Learning," *Nature*, vol. 521, no. 7553, pp. 436–444, 2015.
- [16] F. Marra, D. Gragnaniello, L. Verdoliva, and G. Poggi, "Do GANs Leave Artificial Fingerprints?" in *Proc. IEEE CVPR Workshops on Media Forensics*, pp. 1–8, 2019.
- [17] H. Zhao et al., "Multi-Attentional Deepfake Detection," in *Proc. IEEE/CVF CVPR*, pp. 2185–2194, 2021.
- [18] N. Nakamura, S. Levy, and W. Y. Wang, "r/Fakeddit: A New Multimodal Benchmark Dataset for Fine-grained Fake News Detection," in *Proc. 12th LREC*, pp. 6974–6983, 2020.
- [19] H. Wei, X. Huang, Q. Li, and X. Li, "Towards Multimodal Fake News Detection with Hierarchical Graph Convolutional Network," in *Proc. 29th COLING*, pp. 4164–4175, 2022.
- [20] D. Singh, V. Gupta, and A. Kumaraguru, "MM-COVID: A Multilingual and Multimodal Data Repository for Combating COVID-19 Misinformation," in *Proc. AAAI Workshop on AI for Social Good*, 2021.
- [21] J. Kittler, M. Hatef, R. P. W. Duin, and J. Matas, "On Combining Classifiers," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 20, no. 3, pp. 226–239, 1998.
- [22] D. Cozzolino et al., "ID-Reveal: Identity-aware DeepFake Video Detection," in *Proc. IEEE/CVF ICCV*, pp. 15108–15117, 2021.

- [23] A. Rossler et al., "FaceForensics++: Learning to Detect Manipulated Facial Images," in Proc. IEEE/CVF ICCV, pp. 1–11, 2019.
- [24] P. Korshunov and S. Marcel, "DeepFakes: A New Threat to Face Recognition? Assessment and Detection," arXiv preprint arXiv:1812.08685, 2018.
- [25] K. Shu, D. Mahudeswaran, S. Wang, D. Lee, and H. Liu, "FakeNewsNet: A Data Repository with News Content, Social Context and Spatio-Temporal Information," *Big Data*, vol. 8, no. 3, pp. 171–188, 2020.
- [26] L. Gu et al., "SocialDF: Deepfake Video Detection Benchmark for Social Media," in Proc. IEEE CVPR Workshops on Media Forensics, pp. 1–9, 2022.
- [27] S. Mishra et al., "FACTIFY 2: A Multimodal Fake News and Satire Detection Dataset," in Proc. 37th AAAI Conference on Artificial Intelligence, 2023.
- [28] L. Breiman, "Bagging Predictors," *Machine Learning*, vol. 24, no. 2, pp. 123–140, 1996.
- [29] T. Wolf et al., "Transformers: State-of-the-Art Natural Language Processing," in Proc. 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations, pp. 38–45, 2020.
- [30] M. Sanderson and W. B. Croft, "Deriving Concept Hierarchies from Text," in Proc. 22nd ACM SIGIR Conference on Research and Development in Information Retrieval, pp. 206–213, 1999.